




RAG Docs QA

Ask Gemini

localhost:8501

RAG QA System

Python Documentation AI

Show Sources

Knowledge Base

Build Knowledge Base

LLM Backend

Model

Ollama (Llama 3)

Retrieval Settings

Top-K Chunks

3

RAG-Based QA System

Intelligent Question Answering over Python Documentation (Numpy · Pandas · Matplotlib · Scikit-Learn)

Dense Retrieval (Chroma DB)

Sparse Retrieval (BM25)

LLM Generation

RAGAS Evaluation

Chat

Evaluation

Performance Dashboard

Knowledge Base

System

How do I create a NumPy array from a list?

According to the provided context, you can create a NumPy array from a list using the `np.array()` function. Here's an example:

```
a = np.array([1, 2, 3, 4, 5, 6])
```

This will create a NumPy array with the same elements as the original list.

How can I handle missing values in pandas?

To handle missing values in pandas, you can use the `dropna()` function to drop any rows that have missing data. For example:

```
df1.dropna(how="any")
```

Deploy

⚙️ Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries

💡 Sample Queries

▶ How do I create a NumPy array from a list?

▶ What is the difference between loc and iloc in pandas?

▶ How do I create subplots in Matplotlib?

▶ How does cross-validation work in scikit-learn?

▶ How can I handle missing values in pandas?

Deploy ⋮

🤖 To handle missing values in pandas, you can use the `dropna()` function to drop any rows that have missing data. For example:

```
df1.dropna(how="any")
```

This will return a new DataFrame with all rows that had at least one missing value removed.

Alternatively, you can use the `fillna()` function to fill missing values with a specific value. For example:

```
df1.fillna(value=5)
```

This will replace all missing values in the DataFrame with the value 5.

You can also use the `isna()` function to get a boolean mask where values are nan, and then use this mask to perform operations on your DataFrame.

For example, you could calculate the mean of each column excluding missing values:

```
df.mean()
```

Or you could calculate the mean of each row:

```
df.mean(axis=1)
```

Ask a question about Pandas, Numpy, Sklearn... ➡

⚙️ Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries

💡 Sample Queries

▶ How do I create a NumPy array from a list?

▶ What is the difference between loc and iloc in pandas?

▶ How do I create subplots in Matplotlib?

▶ How does cross-validation work in scikit-learn?

▶ How can I handle missing values in pandas?

🔍 RAG-Based QA System

Intelligent Question Answering over Python Documentation (Numpy · Pandas · Matplotlib · Scikit-Learn)

Dense Retrieval (Chroma DB)

Sparse Retrieval (BM25)

LLM Generation

RAGAS Evaluation

Chat

Evaluation

Performance Dashboard

Knowledge Base

System

📊 System Evaluation

Run a dynamic test suite pulling random knowledge from your built database.

Test Samples

6

Top-K for Eval

4

Library Filter

all

🚀 Run Evaluation

✅ Evaluation complete: 6 unique samples analyzed.

37.0%	100.0%	73.5%	68.5%	63.0%	88.0%
FAITHFULNESS	CTX PRECISION	ANS RELEVANCE	OVERALL	SEMANTIC SIM	CTX RECALL

3

Filter Library

All Libraries

Sample Queries

▶ How do I create a NumPy array from a list?

▶ What is the difference between loc and iloc in pandas?

▶ How do I create subplots in Matplotlib?


▶ How does cross-validation work in scikit-learn?

▶ How can I handle missing values in pandas?


▶ What does numpy broadcasting mean?

▶ How to train a Random Forest classifier?

Dynamic Metric Spider Chart



Performance per Metric



Per-Sample Results

	Question	Library	Faithfulness	Ctx Prec	Ans Rel	Overall	Latency
0	Explain the following concept from sklearn: The algorithm automatically sets the number c	SKLEARN	33.90%	100.00%	71.95%	66.95%	69368ms
1	Explain the following concept from sklearn: Advantages: Interpretability: The unadjusted R	SKLEARN	55.32%	100.00%	82.66%	77.66%	50341ms
2	Explain the following concept from numpy: To get the unique rows, index position, and occ	NUMPY	33.63%	100.00%	71.81%	66.81%	142727ms
3	Explain the following concept from pandas: In [104]: df Out[104]: A B C D E 0	PANDAS	21.43%	100.00%	65.71%	60.71%	166111ms
4	Explain the following concept from numpy: Here, you grabbed a section of your array from	NUMPY	18.75%	100.00%	64.38%	59.38%	100230ms
5	Explain the following concept from sklearn: $\text{FMI} = \frac{\text{TP}}{\sqrt{\text{TP} + 1}}$	SKLEARN	58.75%	100.00%	84.38%	79.38%	150649ms

3

Filter Library

All Libraries

Sample Queries

▶ How do I create a NumPy array from a list?

▶ What is the difference between loc and iloc in pandas?

▶ How do I create subplots in Matplotlib?

▶ How does cross-validation work in scikit-learn?

▶ How can I handle missing values in pandas?

▶ What does numpy broadcasting mean?

▶ How to train a Random Forest classifier?

Chat

Evaluation

Performance Dashboard

Knowledge Base

System

Deploy

Performance Dashboard

8

TOTAL QUERIES

106,854ms

AVG LATENCY

84,097ms

AVG LLM TIME

21,281ms

AVG RETRIEVAL

2,252

TOTAL TOKENS

⚡ Latency over Queries (ms)

Query Index	Total (ms)	LLM (ms)	Retrieval (ms)
0	65000	58000	1000
1	112000	108000	1000
2	70000	53000	18000
3	52000	38000	14000
4	143000	108000	36000
5	165000	125000	42000
6	101000	76000	26000
7	151000	113000	38000

🔵 Token Usage per Query

Query Index	Token Usage
0	650
1	630
2	110
3	85
4	250
5	175
6	160
7	190

🏆 Top Retrieval Score per Query

1

📊 Latency Distribution

RAG QA System

Python Documentation AI

Show Sources

Knowledge Base

Build Knowledge Base

LLM Backend

Model

Ollama (Llama 3)



Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries



Query History

	total_time_ms	llm_time_ms	retrieval_ms	tokens	score	query	type
0	63214.0079	56901.7429	55.9897	656	0.9479	How do I create a NumPy array from a list?	None
1	112188.6322	106303.479	338.048	626	0.9154	How can I handle missing values in pandas?	None
2	69368	52026	17342	111	0.6695	Explain the following concept from sklearn: The algorithm a	System Eval
3	50341	37755.75	12585.25	85	0.7766	Explain the following concept from sklearn: Advantages: Int	System Eval
4	142727	107045.25	35681.75	252	0.6681	Explain the following concept from numpy: To get the uniqu	System Eval
5	166111	124583.25	41527.75	172	0.6071	Explain the following concept from pandas: In [104]: df Out[	System Eval
6	100230	75172.5	25057.5	159	0.5938	Explain the following concept from numpy: Here, you grabb	System Eval
7	150649	112986.75	37662.25	191	0.7938	Explain the following concept from sklearn: $\text{FM} = \sqrt{\text{fr}}$	System Eval

RAG QA System
Python Documentation AI

☒ Show Sources

Knowledge Base

Build Knowledge Base

⌵ LLM Backend

Model

Ollama (Llama 3) ⌵

⚙️ Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries ⌵

RAG-Based QA System
Intelligent Question Answering over Python Documentation (Numpy · Pandas · Matplotlib · Scikit-Learn)

Dense Retrieval (Chroma DB)

Sparse Retrieval (BM25)

LLM Generation

RAGAS Evaluation

Knowledge Base Explorer

228



TOTAL CHUNKS

384

EMBEDDING DIMS

ChromaDB

INDEX TYPE



RAG Docs QA

Ask Gemini

localhost:8501

🔍

☆

🔖

🔍

👤

⋮

 RAG QA System

Python Documentation AI

Show Sources

Knowledge Base

Build Knowledge Base

⌵ LLM Backend

Model

Ollama (Llama 3)

⚙️ Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries

19.3%

3.07%

50

0

guide

tutorial

Deploy

 Browse Chunks

Filter by Library

All

Search in chunks

e.g. array creation...

Showing 10 of 228 matching chunks

> [PANDAS] 10 minutes to pandas — pandas 3.0.2 documentation ...

> [PANDAS] API reference Development Release notes ...

> [PANDAS] User Guide 10 minutes to pandas 10 minutes to panda...

> [PANDAS] In [4]: s Out[4]: 0 1.0 1 3.0 2 5.0 3 NaN 4 ...

> [PANDAS] Creating a DataFrame by passing a dictionary of objects wher...

> [PANDAS] If you're using IPython, tab completion for column names (as...

> [PANDAS] Viewing data# See the Essential basic functionality....

> [PANDAS] In [16]: df.columns Out[16]: Index(['A', 'B', 'C', 'D'], dtype=...

> [PANDAS] In [19]: df2.to_numpy() Out[19]: array([[1.0, Timestamp('20...

> [PANDAS] Transposing your data: In [21]: df.T Out[21]: 2013-01-01...

RAG Docs QA

Ask Gemini

localhost:8501

RAG QA System

Python Documentation AI

Show Sources

Knowledge Base

Build Knowledge Base

LLM Backend

Model

Ollama (Llama 3)

Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries

RAG-Based QA System

Intelligent Question Answering over Python Documentation (Numpy · Pandas · Matplotlib · Scikit-Learn)

Dense Retrieval (Chroma DB)

Sparse Retrieval (BM25)

LLM Generation

RAGAS Evaluation

Chat

Evaluation

Performance Dashboard

Knowledge Base

System

System Configuration

Service Status

Ollama Backend: **CONNECTED**

Current Vector DB: **ChromaDB**

Embedding Model: **all-MiniLM-L6-v2**

Hardware Allocation

RAM Usage: 81.5% of 7.7 GB

Compute Device: CPU (Thread Optimized)

Quick Start Commands

Select Operating System:

Windows

Linux/Mac

```
# 1. Install dependencies
pip install -r requirements.txt
```

RAG QA System
Python Documentation AI

☒ Show Sources

Knowledge Base

Build Knowledge Base

LLM Backend

Model

Ollama (Llama 3)

Retrieval Settings

Top-K Chunks

3

Filter Library

All Libraries

Quick Start Commands

Select Operating System:

☒ Windows ☐ Linux/Mac

```
# 1. Install dependencies
pip install -r requirements.txt

# 2. Install Ollama (Download Installer)
# Visit: https://ollama.com/download/windows

# 3. Pull a model (in PowerShell/CMD)
ollama pull llama3

# 4. Start the RAG App
streamlit run app.py
```

Project Structure

